

File Systems

Operating Systems

File Systems

- What is a **file**? 
- Array of **persistent** bytes that can be read/written

- **File system**




- Collection of files
- Also, part of OS that manages those files
- Files need *names* to access correct file

inode

- **Kernel** structure for file-system object
 - e.g., file or directory


- Stores **metadata**
 - Type (file / directory)
 - UID, GID (owner / group)
 - rwx (permissions)
 - Size (in bytes)
 - Data blocks
 - Times (access / create)
 - Links count (# paths)
 - ...



inode

- Each file has exactly one **inode** number
- Inodes are unique
 - At a given time
 - **Within the file system**
 - Stored in **inode table**
- See inodes via `ls -i`

inode



File API




- 1st attempt:
 - **read**(int inode, void *buf, size_t nbyte)
 - **write**(int inode, void *buf, size_t nbyte)
 - **seek**(int inode, off_t offset)
- Cons?
 - Names hard to remember
 - No organization or meaning to inode numbers
 - Offset across multiple processes?

File Names

- **File Names**
 - String names are friendlier
 - File system still interacts with *inode* number
- Store *name-to-inode* mapping in predetermined file
 - Typically, inode **2**


File Names



File Paths

- Let's generalize paths!
 - **Directory tree** instead of single root
 - Only **full paths** need to be unique
 - /usr/myuser/file.txt
 - /tmp/file.txt
- Store *file-to-inode* mapping for each **directory**

File Paths



File API




- 2nd attempt:
 - `read(char *path, void *buf,
off_t offset, size_t nbyte)`
 - `write(char *path, void *buf,
off_t offset, size_t nbyte)`
- Cons?
 - Expensive traversal
 - Goal: traverse *once*

File Descriptors




- **File Descriptor (fd)**
 - Do expensive traversal once
 - Store inode in *descriptor* object
 - Do operations & track offset via descriptor
- Each process:
 - File descriptor table: pointers to open file descriptors
 - Integers for file I/O indexed into table
 - stdin: 0, stdout: 1, stderr: 2

File Descriptors



File Descriptors



```
int fd1 = open("file.txt"); // returns 3
read(fd1, buf, 12);
int fd2 = open("file.txt"); // returns 4
int fd3 = dup(fd2);           // returns 5
```

File API



- 3rd attempt:
 - `int open(char *path, int flag, mode_t mode)`
 - `read(int fd, void *buf, size_t nbyte)`
 - `write(int fd, void *buf, size_t nbyte)`
 - `close(int fd)`
- String names
- Hierarchical
- Traverse once
- Different offsets precisely defined

Hard Links

- No system call to delete files!


Two paths/names for the same file?

Hello	1415
readme	1390
file1.txt	4200
file2.txt	505
file3.txt	4200
a.out	220
testfile	230


- File paths are called **hard links**
 - Across *same* or *different* directories
 - inode tracks # of hard links
 - File is **garbage collected** when there are no references

Symbolic Links

- **Symbolic links**
 - Special file, points to another *path name*
 - Like Windows shortcuts



Links



File System Implementation


- File system: pure software
 - Many different file systems exist
- Start with case study: **VSFS**
 - Simplified version of typical UNIX file system

Very Simple File System

- **Data structures**
 - What type of on-disk structures?
- **Access methods**
 - How to map calls (`open()`, `read()`, `write()`, etc.)?
 - Read which structures during which calls?


Overall Organization

- Divide disk into **blocks**
 - Addressed 0 to $N-1$, commonly-used size: 4KB




Overall Organization

- Reserve **data region** for user data
 - e.g., fixed portion: 54 of 63 blocks



Overall Organization

- Reserve blocks for file **metadata** (inode)
 - Typically 256 bytes per inode
 - 4KB disk block → 16 inodes per inode block



Overall Organization

- **Inode**

- Type (file / directory)
- UID (owner)
- rwx (permissions)
- Size (in bytes)
- Data blocks
- Times (access / create)
- Links count (# paths)
- ...

inode 16	inode 17	inode 18	inode 19
inode 20	inode 21	inode 22	inode 23
inode 24	inode 25	inode 26	inode 27
inode 28	inode 29	inode 30	inode 31

Allocation

- Where is file data located?
- Need a way to allocate data blocks to files
 - Store information in inode
 - Creates external fragmentation?
 - Can the file grow?
 - What is the performance?
 - Metadata overhead?

Allocation


- Contiguous allocation
 - Metadata: starting block and file size (**single extent**)
 - OS allocates by finding sufficient free space
 - Must predict future size of file to reserve enough space



- Fragmentation: Bad external fragmentation
- Grow file: Not without moving
- Sequential access: Excellent performance
- Random access: Simple calculation
- Metadata: Very little overhead (“wasted space”)

Allocation


- **Linked-list**
 - Inode pointer to 1st block, end of data → another pointer



- Fragmentation: No external fragmentation
- Grow file: Can grow easily
- Sequential access: Depends on data layout
- Random access: Very poor (traverse all blocks)
- Metadata: Pointer per block wasted

Allocation


- File-Allocation Table (**FAT**)
 - Linked-list information in on-disk FAT table
 - Directory entries instead of inodes (hard links impossible)



- Read from two disk locations for every data read
- Optimize: cache FAT in main memory
 - Greatly improves random access
 - What to cache? Does not scale with file system

Allocation


- Indexed allocation
 - Fixed-size blocks for each file, as described earlier



- Fragmentation: No external fragmentation
- Grow file: Can grow easily
- Sequential access: Depends on data layout
- Random access: Good
- Metadata: Wasted space for pointers (most files are small)


Allocation

- Where are the data blocks?
 - Inode contains one or more *direct* pointers
 - File size is limited




Allocation

- **Indirect pointers**
 - Point to a block that contains more pointers
 - Indirect blocks are stored in regular data blocks




Allocation

- Optimize for small files
 - Fixed number of direct pointers, single indirect pointer
 - Also: double indirect, triple indirect, etc.



Multi-Level Index


- Double indirect: points to block of indirect pointers
- Triple indirect: points to block of double indirect pointers



- Example:
 - Block size 4KB, 4-byte pointers
 - 12 direct pointers, 1 single and 1 double indirect pointers
 - Can accommodate 4GB file
 - $((12 + 1024 + 1024^2) \cdot 4KB)$

Allocation

- Extents
- Metadata: small array (2-6) of extents




- Fragmentation: Helps external fragmentation
- Grow file: Can grow (until out of extents)
- Sequential access: Excellent
- Random access: Simple
- Metadata: Small overhead

Multi-Level Index

- Many file systems use multi-level index
 - Linux ext2 and ext3, NetApp's WAFL, Unix file system
 - SGI XFS and Linux ext4 use **extents**
- **Extents**
 - Disk pointer plus length
 - Avoids large metadata per file


Allocation

- How do we find free data blocks or free inodes?
- **Free list**
- **Bitmap:** for data region and inode table



Superblock

- On mount, OS reads **superblock** to initialize
 - Information about file system



Access Paths

- Open file /foo/bar (assume it exists)
- Traverse pathname to locate desired inode
 - Begin at root: well-known, usually inode 2
 - Read block that contains inode 2
 - Look inside It – read data block to find inode number of foo
 - Read inode of foo to find data block of foo
 - Read data blocks of foo to find inode of bar
 - Read inode of bar
- 5 reads just to open a file!

Access Paths

- Write to /foo/bar (assume it is open)
 - Read data bitmap
 - Write updated data bitmap (newly-allocated block to use)
 - Read the inode of bar
 - Write updated inode with new block location
 - Write the actual block itself
- Read /foo/bar (assume it is open)
 - Read bar inode to find data block
 - Read data block
 - Write bar inode – update access time

Access Paths

- Create file /foo/bar (doesn't exist)
 - Read inode bitmap, write updated bitmap with allocated inode
 - Read block of found inode and write inode itself for bar
- Find data block of foo
 - Read block that contains inode 2 (root)
 - Read data block to find inode number of foo
 - Read inode of foo to find data block of foo
 - Read and write data block of foo to add bar entry to directory
 - Write update inode of foo (date change)
- Directory foo needs to grow? Additional I/O

Caching and Buffering

- Simple operations, huge number of I/Os
- Long pathnames can lead to hundreds of reads
- Solution? Use memory (RAM) to cache important blocks
 - **Write buffering**
 - Batch updates for less I/O (several updates to inode bitmap)
 - Buffer writes in memory, **schedule** subsequent I/Os
 - **Avoid** writes, e.g., file created and then deleted

Very Simple File System

- Described a very simple FS
 - Basic on-disk structures
 - Basic operations
 - Poor performance
- In practice:
 - Allocate **efficiently** to obtain good performance
 - Handle **crashes**


Crash Consistency

- File system data structures must **persist**
- Major challenge: update structures despite **power loss** or **system crash**
- **Crash-consistency problem**
- On-disk structures left in an **inconsistent state**

Crashes may occur at arbitrary times.
How do we ensure the file system remains valid?


Example

- Append 4KB to an existing file
 - Open file → lseek () → issue 4KB write → close file
- Simple file system:
 - **Inode bitmap, data bitmap, inodes, data blocks**




Example

- Write three blocks:
 - Updated data bitmap
 - Updated inode
 - New data block




Crash Scenarios

- Crash after a single write:
 - Updated data bitmap




Crash Scenarios

- Crash after a single write:
 - Updated inode




Crash Scenarios

- Crash after a single write:
 - New data block




Crash Scenarios

- Crash after *two writes*:
- Updated inode and data bitmap




Crash Scenarios

- Crash after *two writes*:
- Updated inode and new data block



Crash Scenarios

- Crash after *two writes*:
- Updated data bitmap and new data block



Crash Scenarios

- Goal: move *fs* from consistent state to another **atomically**
 - Don't care about data loss, only consistency
 - But disk only commits **one** write at a time
 - Crash or power loss may occur
- **Crash-consistency problem**

Solution #1

- Let inconsistencies happen, fix them later
- File System Checker (`fsck`)
 - UNIX tool for finding and repairing inconsistencies
 - Can't fix all problems
 - e.g., file system looks consistent but inode points to garbage data
 - Only goal: keep file system *metadata* consistent

Solution #1

- **Free blocks**
 - Scan **inodes** to produce correct version of data bitmap
- **Superblock and Inode state**
 - Make sanity checks on file system & allocated blocks
 - Integrity checks on directory contents
 - Check each **inode** for corruption
 - Verify link count of each allocated **inode**
- **Duplicates**
 - Check for duplicate points to same data block
 - Clear bad inode or duplicate block

Solution #1

- Building `fsck` requires intricate knowledge of file system
- Too slow
 - Scan entire disk to find allocated blocks
 - Read entire directory tree
 - Many minutes or *hours*
- Previous example:
 - Just three blocks are written to disk
 - Scan entire disk for problems

Solution #2

- **Journaling (write-ahead logging)**
 - Add work during updates → reduce work during recovery
 - First, write a *note* describing operation
 - If crashed, look at *notes* and try again
 - Know exactly what to fix instead of scanning entire disk




Fsync

- File systems keep newly written data in memory
 - Write buffering improves performance
- What if system crashes?
 - `fsync(int fd)`
 - Forces buffers to flush to disk
 - Usually tells disk to flush its write cache as well

Journaling

- New key structure: **journal** itself
 - Occupies small amount of space within partition
 - ...or another device



Journaling

- **Checkpoint:** transaction safely in **journal**
 - Ready to overwrite structures in file system



- Initial sequence:
 1. **Journal write**
 - Write the transaction
 - Wait for writes to complete
 2. **Checkpoint**
 - Write data and metadata to final locations

Journaling

- Crash during **journal** write?



- Issue writes in two steps



Journaling

- Updated protocol:

1. Journal write

- Write transaction contents (TxB , *metadata*, *data*)
- Wait for writes to complete

2. Journal commit

- Write transaction commit block (TxE)
- Wait for write to complete → transaction is **committed**

3. Checkpoint

- Write data and metadata to final locations

Recovery


- Crash before transaction is written to log?
 - Pending update is skipped
- Crash after commit, before checkpoint completes?
 - Can **recover** the update
 - On boot, scan log for committed transactions to **replay**
 - Called **redo logging**

Finite Log

- Log of finite size:
 - Larger log → longer recovery
 - Log full → no more transactions
- Solution? **Circular log**
 - Free log space used by the transaction
 - Some time after a checkpoint

Finite Log

- **Journal superblock**
 - Marks oldest and newest non-checkpointed transactions




Finite Log

- Another step in our protocol
 - 1. Journal write**
 - 2. Journal commit**
 - 3. Checkpoint**
 - 4. Free**
 - Some time later, mark transaction as free in journal
 - By updating the journal superblock


Metadata Journaling

- Problem: writing every data block twice
 - Commit to log
 - Checkpoint to disk



Metadata Journaling

- **Ordered journaling:** user data is not written to journal
 - Write data block to file system proper



- Write data *after* transaction → may point to garbage data
- Write data *before* transaction → avoids problem

Journaling

- Final protocol:

1. Data write

- Write data to final location, wait for completion

2. Journal write

- Write transaction contents (TxB and metadata), wait for completion

3. Journal commit

- Write transaction commit block (**TxE**)
- Wait for write to complete → transaction is **committed**

4. Checkpoint

- Write metadata to final locations

5. Free

- Later, mark transaction as free in journal superblock